

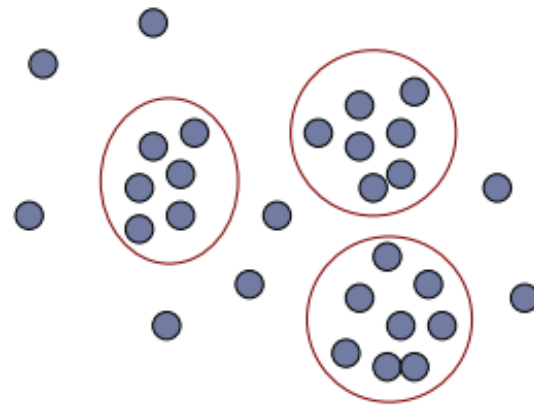
Cluster Analysis

Jiaying Liu G45268292

Hao Wu G36059463

What is Cluster Analysis

- Unsupervised learning (i.e., class label is unknown)
- Group data to form new categories (i.e., clusters), e.g., cluster houses to find distribution patterns
- Principle: Maximizing intra-class similarity & minimizing interclass similarity



- Typical Applications: “Market basket data”, “Amazon recommends products”, “biology clustering” etc.

Supervised VS. Unsupervised Learning

- Supervised: Dataset with pre-defined classification that we can use this to supervised learn the dataset and create the model
- Example: discriminant analysis, logistic regression
- Unsupervised: The classification categories are unknown, but features are observed that relate to the unobserved categories.
- Example: cluster method

Measuring Dissimilarity Between Observations

- For **categorical data**, vectors of observations on p binary variables.
- $P=a+b+c+d$
- Similarity (matching index): $\frac{a + d}{a + b + c + d}$
- Dissimilarity: $\frac{b + c}{a + b + c + d}$

cross classification of two observations

on p binary variables		
observation h	Observation i	
	1	0
1	a	b
0	c	d

Measuring Dissimilarity Between Observations (cont.)

- In special applications, value 1 is more common (or meaningful) than value 0.
- Jaccard index: $\frac{a}{a+b+c}$
- Dissimilarity: $\frac{b+c}{a+b+c}$

cross classification of two observations on p binary variables		
observation h	Observation i	
	1	0
1	a	b
0	c	d

Measuring Dissimilarity Between Observations (cont.)

- For **continuous data**, we usually calculate the distance between sample points.
- Treat data as point (or vector) in the dimensional space.
- Smaller distance, larger similarity.
- **Minkowski distance:**
 - For $p = [p_1, p_2, \dots, p_m]$ and $q = [q_1, q_2, \dots, q_m]$
 - $d_x(p, q) = (\sum_{i=1}^m |p_i - q_i|^x)^{1/x}$, ($x > 0$)

Measuring Dissimilarity Between Observations (cont.)

- Minkowski distance: $d_x(p, q) = (\sum_{i=1}^m |p_i - q_i|^x)^{1/x}$, ($x > 0$)

- $x=1$,

- $d_1(p, q) = \sum_{i=1}^m |p_i - q_i|$

Hamming distance

- $x=2$,

- $d_2(p, q) = \sqrt{\sum_{i=1}^m |p_i - q_i|^2}$

Euclidean distance

- $x = \infty$,

- $d_\infty(p, q) = \max_{1 \leq i \leq m} |p_i - q_i|$

Chebyshev distance

May be effected by the unit of measurement, **require normalization before use!**

Measuring Dissimilarity Between Observations (cont.)

- **Canberra distance:** $d_{canb}(p, q) = \sum_{i=1}^m \frac{|p_i - q_i|}{|p_i| + |q_i|}$
- Overcome the influence of measurement unit
- When coordinates of two points are both close to 0, Canberra distance is sensitive with tiny changes.
- Besides distance, we can also use coefficient (like correlation coefficient) to measure similarity.

Clustering Categories

➤ Partitioning Methods

- Construct K partitions of the data
 - ❖ K-mean Clustering
 - ❖ K-medoid Clustering

➤ Hierarchical Methods

- Creates a hierarchical decomposition of the data
 - ❖ Bottom – up / agglomerative
 - ❖ Top-down / divisive

Partitioning Methods: The Principle

➤ Given

- A data set of n objects
- K the number of clusters to form

➤ Organize the objects into K partitions ($k \leq n$) where each partition represents a cluster

➤ The clusters are formed to optimize an objective partitioning criterion

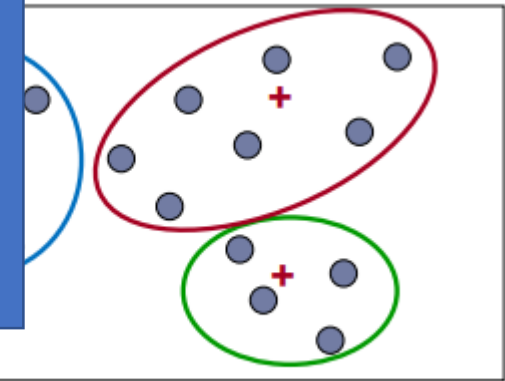
- Objects within a cluster are similar
- Objects of different clusters are dissimilar

Partitions---K-mean clustering

- Idea: Find the partition that minimizes the total dissimilarity

Actually it is not easy to minimize the square-error function, to find the optimal solution, we need to find all the possible cluster. So k mean use the greedy algorithm to find the literately approximant solution.

smallest total



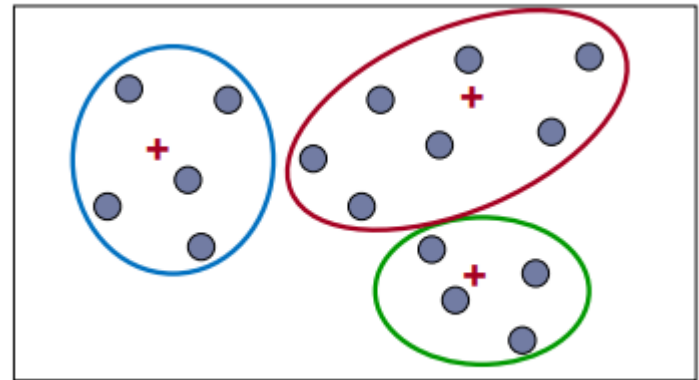
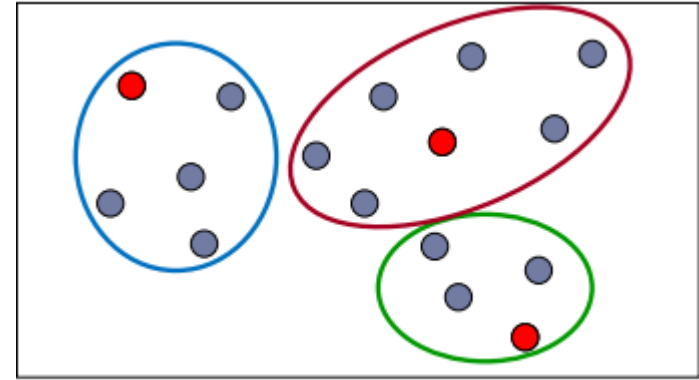
- The algorithm iteratively updates the partition that minimizes the square-error function

$$E = \sum_{i=1}^k \sum_{p \in C_i} (p - m_i)^2$$

- It works well when the clusters are compact clouds that are rather well separated from one another

Basic Process of K-mean Clustering

- 1. Arbitrary choose k objects from D as in initial cluster centers
- 2. Assign each object to the most similar cluster based on the mean value of the objects in the cluster



Process of K-mean clustering

- 3. U

Question:

1. How to find the initial value?
2. How to cluster and update ?
3. How to decide that the algorithm is converged(stop sign)?

- 4.If



- **Input** $D = \{x_1, x_2, x_3, \dots, x_n\}$; # of Cluster k .
- { Randomly choose k points as initial $\mu = \{\mu_1, \mu_2, \mu_3, \dots, \mu_k\}$

Repeat

let $C_i = \emptyset$ ($1 \leq i \leq k$)

for $j = 1, 2, \dots, n$ **do**

calculate each points x_j distance with each μ_i : $d_{ij} = \|x_j - \mu_i\|_2$

find the minimum distance and set it to that cluster

end for

for $i = 1, 2, \dots, k$ **do**

calculate the new $\mu'_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$;

if $\mu'_i \neq \mu_i$ **then**

update $\mu_i = \mu'_i$

else

keep the μ_i

end if

end for

until μ_i converge

Output cluster = $C = \{C_1, C_2, C_3, \dots, C_k\}$

K-Mean Properties

Advantages

- K-means is relatively scalable and efficient in processing large data sets
- The computational complexity of the algorithm is $O(nkt)$

Disadvantage

- it is applicable only when the mean is defined (Cannot apply for categorical data)
- Users need to specify k
- K-means is not suitable for discovering clusters with nonconvex shapes or clusters of very different size
- It is sensitive to noise and outlier data points (can influence the mean value)

Optimize the algorithm

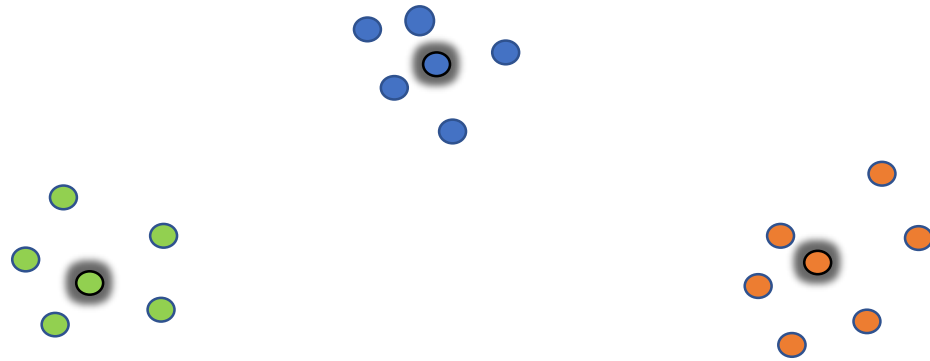
- How to find the K?
 - Plots the value of the clustering criterion (cost function) against k , looking for a natural break point where this changes substantially
 - Plot a summary against k such as the probability that the dissimilarity for a with-in cluster pair is smaller than the dissimilarity for a between-cluster pair.

- What initial value can make algorithm converge more quickly?
 - Even we say we can randomly choose k point and since the algorithm will update through each iteration. In the end, the mean will move and eventually go to the right place;
 - But image: what is it take a lot of time? what if it vibrate? What if the point we choose is outlier????
 - First estimate the center of whole dataset and create k point which are the value of scale multiple random normal value like:

```
group_center = mean(data_set);  
group_range = range(data_set);  
centers = (randn(K,dim).* repmat(group_range,K,1)./3+ repmat(group_center,K,1));
```

Partitions---K-medoid clustering (PAM)

- Medoid: The medoid of a cluster is the truly observation with smallest total dissimilarity to the other points in the cluster



- The goal of k-medoid clustering is to minimize the sum of the dissimilarities between each object and its corresponding reference point (medoid)

$$E = \sum_{i=1}^k \sum_{p \in C_i} |p - o_i|$$

Process of K-medoid clustering

Step1 : Initialize randomly select k of the n data points as the medoids

Step2 : Associate each data point to the closest medoid.

Step3 : for each medoid m

 if each non-medoid data point o

 swap m and o and compute the total cost of the configuration

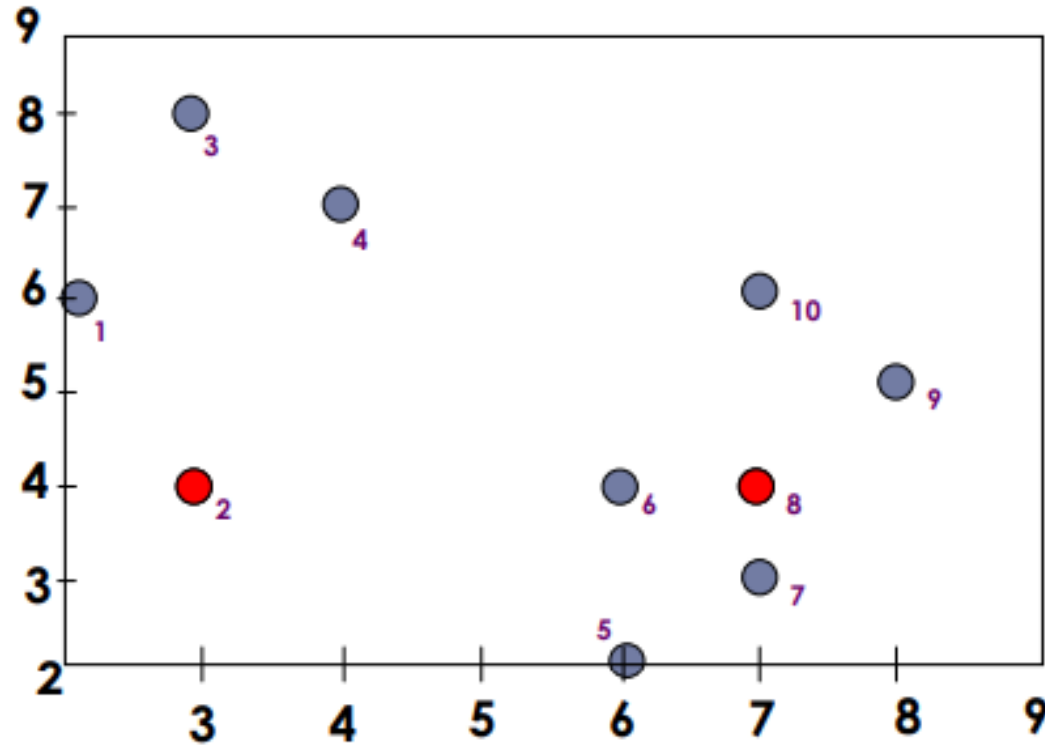
Step4 : Select the configuration with the lowest cost.

Step5 : Repeat steps 2 to 5 until there is no change in the medoid.

Example of k-medoid clustering

Data objects

	A_1	A_2
O_1	2	6
O_2	3	4
O_3	3	8
O_4	4	7
O_5	6	2
O_6	6	4
O_7	7	3
O_8	7	4
O_9	8	5
O_{10}	7	6

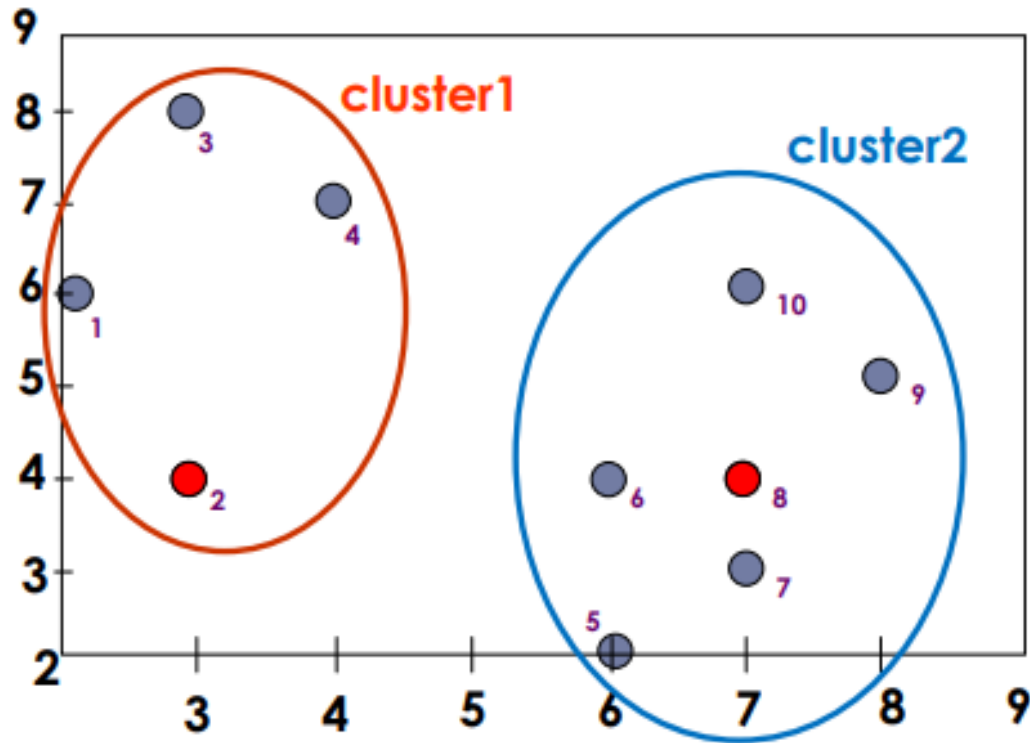


Goal: create two cluster

Choose randomly two medoids

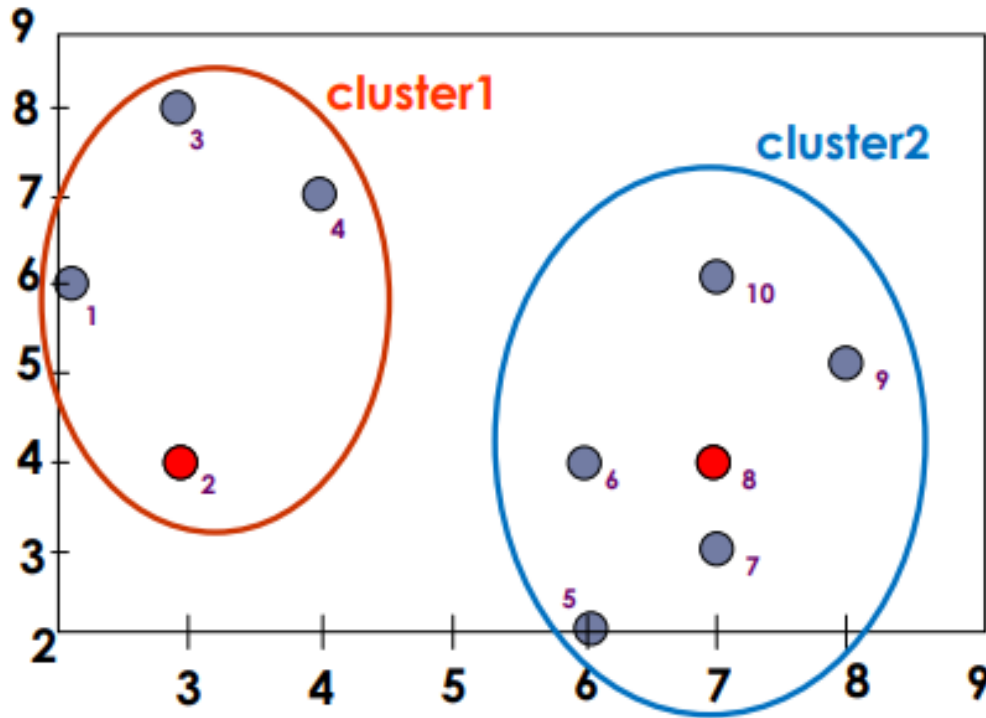
$O_2(3,4)$, $O_8(7,4)$

Example of k-medoid clustering



- Assign each object to the closest representative object
- Using distance to form the following clusters
- Cluster1= {O1,O2,O3,O4}
- Cluster2={O5,O6,O7,O8,O9,O10)

Example of k-medoid clustering

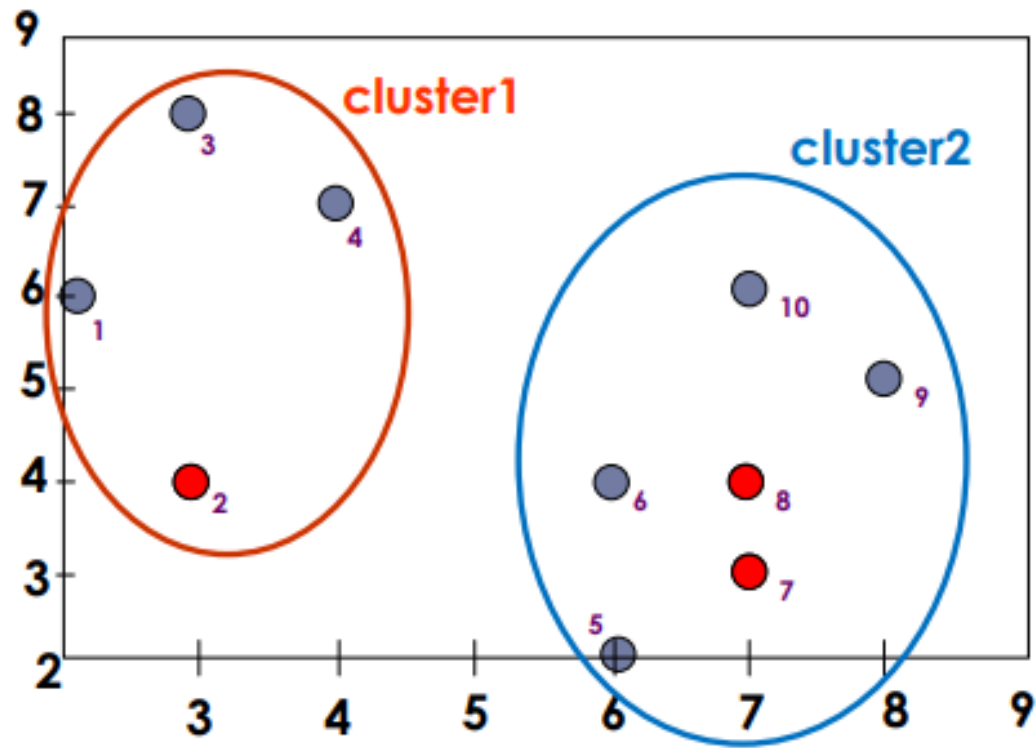


- Compute the absolute error function [for the set of Medoids (O2,O8)]

$$E = \sum_{i=1}^k \sum_{p \in C_i} |p - o_i| = |o_1 - o_2| + |o_3 - o_2| + |o_4 - o_2| \\ + |o_5 - o_8| + |o_6 - o_8| + |o_7 - o_8| + |o_9 - o_8| + |o_{10} - o_8|$$

$$E = (3 + 4 + 4) + (3 + 1 + 1 + 2 + 2) = 20$$

Example of k-medoid clustering



Choose a random object O7

Swap O8 and O7

Compute the absolute error function [for the set of Medoids (O2,O7)]

$$E = (3 + 4 + 4) + (2 + 2 + 1 + 3 + 3) = 22$$

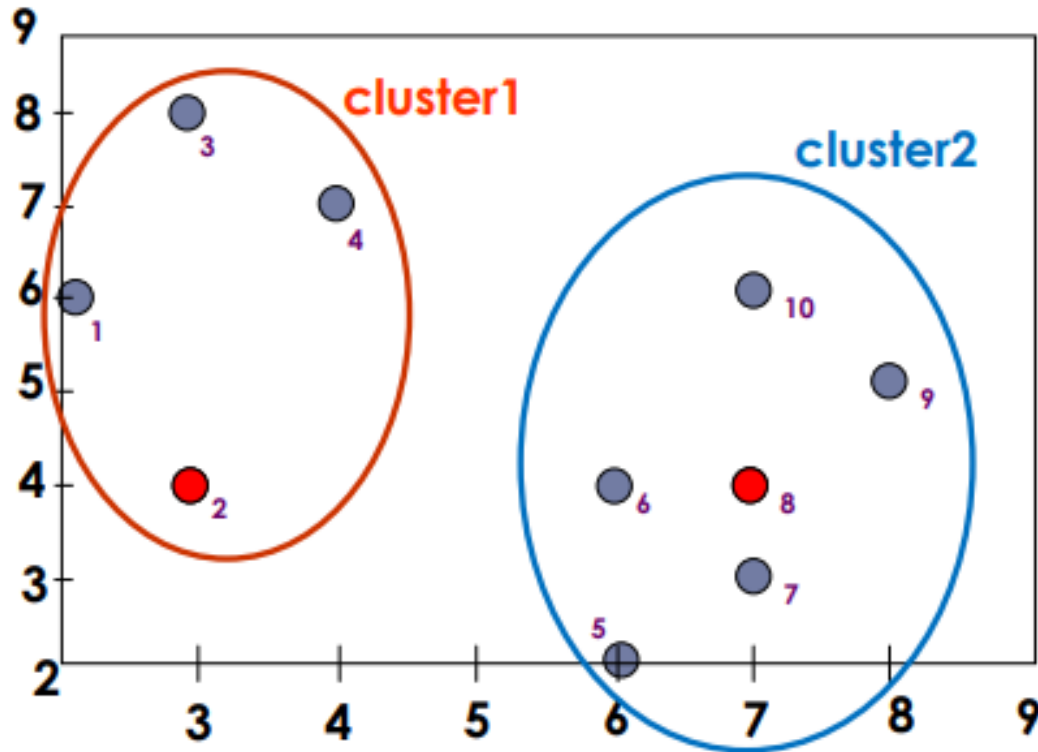
Compute the cost function

Absolute error [for O2,O7] – Absolute error [O2,O8]

$$S = 22 - 20$$

$S > 0$ It is a bad idea to replace O8 by O7

Example of k-medoid clustering



- Since there is no change in the medoid set, the algorithm ends here. Hence the clusters obtained finally are
- Cluster1= {O1,O2,O3,O4}
- Cluster2={O5,O6,O7,O8,O9,O10}

K-medoids properties(k-medoids vs. k-means)

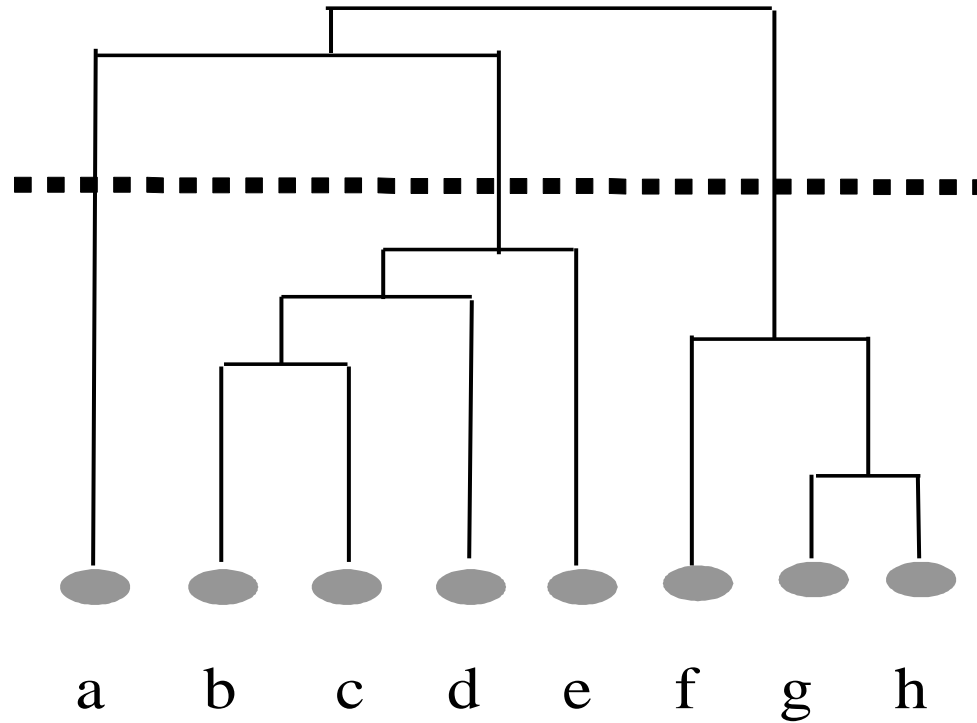
- Advantages
 - K-Medoids method is more robust than k-Means in the presence of noise and outliers
- Disadvantages
 - K-Medoids is more cost than the k-Means method
 - Like k-means, k-medoids requires the user to specify k
 - It does not scale well for large data sets
 - (CLARA[Clustering Large Application], CLARANS[Clustering Large Application based upon Randomized Search])

Hierarchies clustering

- HC provides graphical illustration of relationships between the data in the form of a dendrogram (binary tree).
- Two approaches: **agglomerative** & **divisive**
- Agglomerative / bottom-up method starts with each object in the data forming its own cluster, and then successively merges the clusters until one large cluster is formed, which encompasses the entire dataset.
- Divisive / top-down method starts by considering the entire data as one cluster and then splits up the cluster(s) until each object forms its own cluster.

Dendrogram

Top-down /
divisive



Bottom-up /
agglomerative




{a}
{b,c,d,e}
{f,g,h}

Procedure of Agglomerative Clustering

Given: a data set and the distance function

1. start with “N” clusters by assigning each pattern to a separate cluster
2. proceed with this initial configuration of the clusters and merge the clusters that are the closest. In other words, if S and T are the two clusters being recognized as the closest, form a single cluster {S, T} and reduce the number of clusters by one
- 3 repeat step 2 until a minimal number of the clusters has been reached.

Set a threshold. If the smallest distance less than threshold, stop iteration.



Result: clusters of data (partition)

Distance Between Clusters

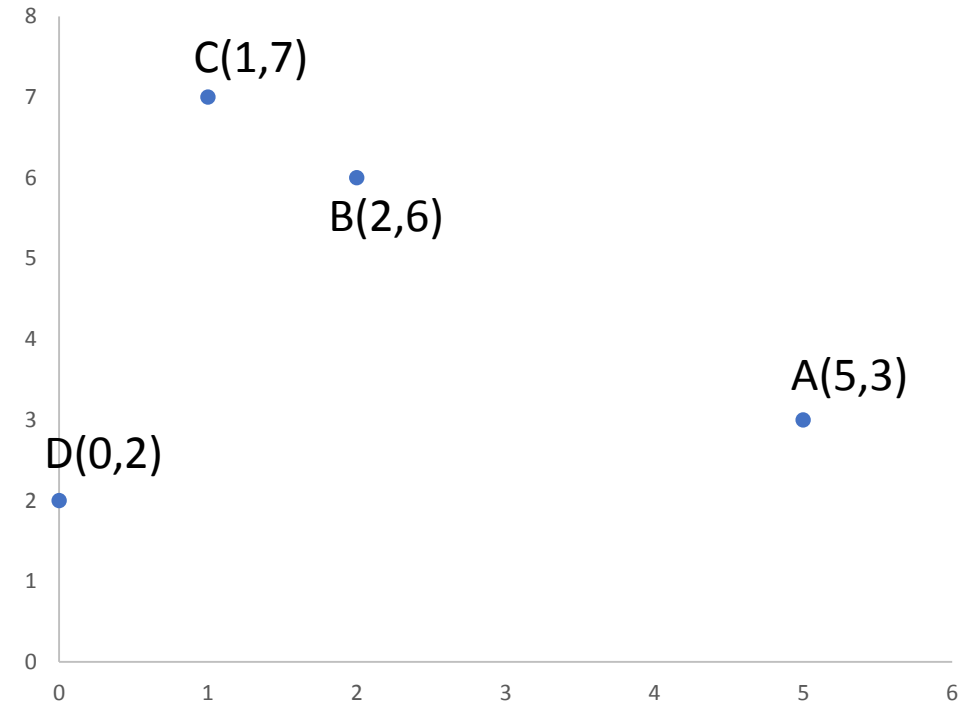
Single linkage method : $\|T-S\| = \min_{\substack{x \in T \\ y \in S}} \|x - y\|$

complete linkage : $\|T-S\| = \max_{\substack{x \in T \\ y \in S}} \|x - y\|$

average linkage : $\|T-S\| = \frac{1}{card(S)card(T)} \sum_{\substack{x \in T \\ y \in S}} \|x - y\|$

Example

- Suppose 4 data: A(5,3), B(2,6), C(1,7), D(0,2)
- We want to divide data into two clusters.
- Use agglomerative clustering.
- Measure dissimilarity between observations with Euclidean distance.
- Measure distance between clusters with average linkage.



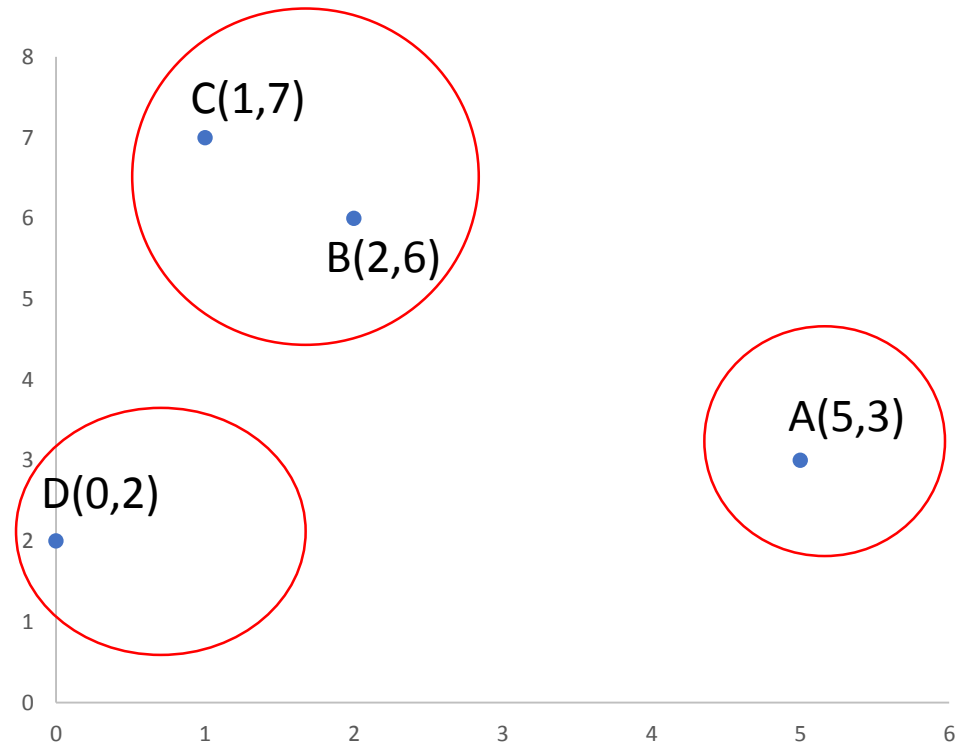
Example

- At initial stage: treat each observation as a single cluster.

- Calculate dissimilarity:

	A	B	C	D
A		4.24	5.66	5.10
B	4.24		1.41	4.47
C	5.66	1.41		5.10
D	5.10	4.47	5.10	

- Combine the clusters with smallest distance as one cluster.
- Now 3 clusters: {A}, {B,C}, {D}



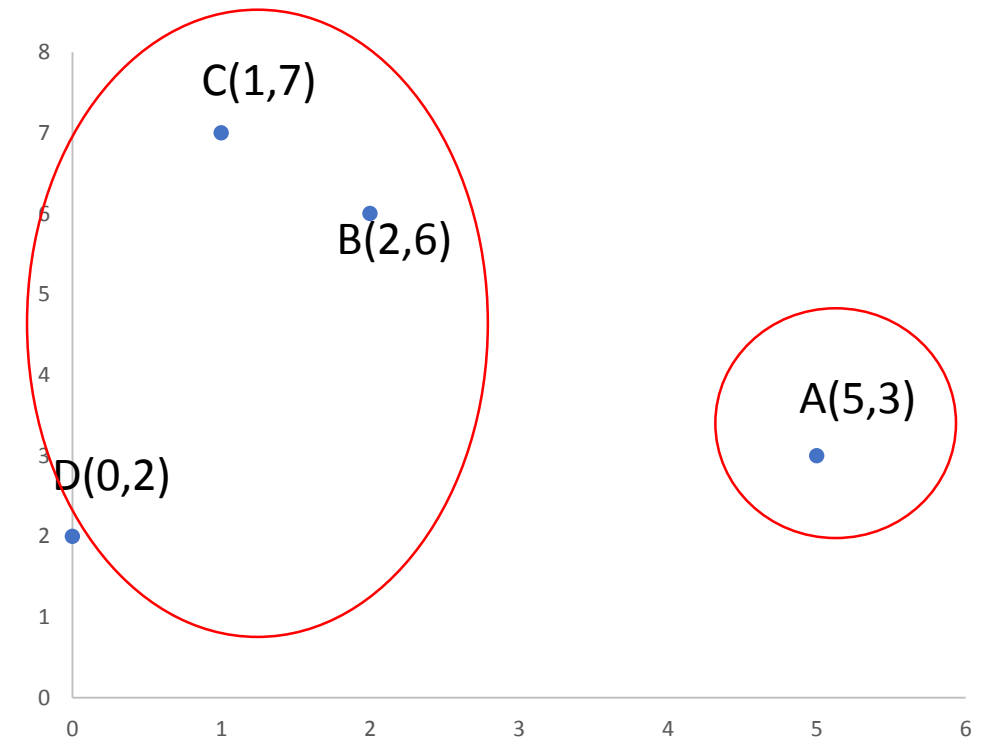
Example

- Calculate cluster distance:

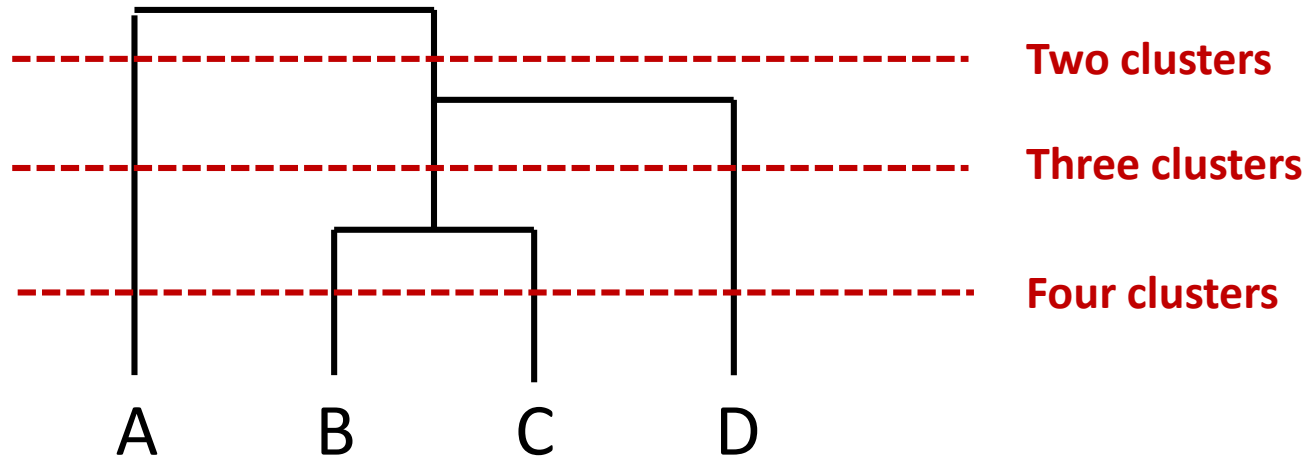
	A	B&C	D
A			
B&C	4.95		4.79
D	5.10	4.79	

- Combine the clusters with smallest distance as one cluster.
- Now 2 clusters: {A}, {B,C,D}

$$\begin{aligned} &\text{Average linkage between } \{A\} \text{ and } \{B,C\} \\ &= \frac{1}{2} [\sqrt{(A-B)^2} + \sqrt{(A-C)^2}] \end{aligned}$$



Example



Reference

- <https://cran.r-project.org/web/packages/kmed/vignettes/kmedoid.html>
- <https://www.mathworks.com/help/stats/kmedoids.html>
- <http://www.sthda.com/english/wiki/print.php?id=236>
- <https://www.datanovia.com/en/lessons/clara-in-r-clustering-large-applications/>
- https://en.wikibooks.org/wiki/Data_Mining_Algorithms_In_R/Clustering/C_LARA
- <http://www.sthda.com/english/wiki/print.php?id=239>
- <https://cran.r-project.org/web/packages/cluster/cluster.pdf>
- <https://cran.r-project.org/web/packages/kmed/kmed.pdf>
- <https://cran.r-project.org/web/packages/kmed/vignettes/kmedoid.html>